Combining Truth Discovery and Open Information Extraction with Active Ensembling

Mouhamadou Lamine BA
Qatar Computing Research Institute
Tornado Tower, West Bay
Doha, Qatar
mlba@qf.org.qa

Laure Berti-Equille
Qatar Computing Research Institute
Tornado Tower, West Bay
Doha ; Qatar
lberti.qf.org.qa

ABSTRACT

Web search engines or open information extraction systems usually reply to users' queries with a set of candidate answers that are often conflicting because they are claimed by multiple information sources. In this context, estimating information veracity is difficult for the users especially when they have no prior knowledge about the trustworthiness of the sources. In this demonstration paper, we showcase a system that supports event/entity and relation extraction based on keyword-search from the Web, processes the conflicting outputs, combines multiple truth finding algorithms with active learning to provide the most likely true answers and determine the most trustworthy sources.

1. INTRODUCTION

[**Lamine: Page allocation**]
- 1.25 pages –> abstract + introduction
- 1.5 pages –> Open information extraction + Active Ensembling for Truth Discovery
- 1 pages –> Demonstration System + Scenario
- 0.25 pages –> References

2. OPEN INFORMATION EXTRACTION

- décrire le type d'information auquel on s'intéresse par exemple "factoid claim"
- decrire le systeme sur lequel on se base
- décrire comment on transforme l'output de OpenIE
- donner qq exemples

We are seeking to demonstrate in this paper the usefulness of truth discovery on large sets of *"factoid" claims* about real-world facts which are obtained when querying *open information extraction* (OpenIE) systems. These claims are usually extracted by information extractors from unreliable and conflicting Web sources. A "factoid" claim, e.g., Barack Obama was born in Kenya, often refers to a piece of information that is accepted as true, without any prior verification, because of its frequent redundancy over numerous sources.

Multiple conflicting Web claims of such a type related to a specific real life fact are simultaneously available in practice. We draw, therefore, our attention on those conflicting claims returned by a common open Web information extraction system, *TextRunner* in our special case, as possible answers to users' queries. Concretely speaking, given a user input query (or key phrase), our main goal is to consider and improve the result of TextRunner[1] by running truth discovery in order to provide to the user the most reliable information. We next briefly present TextRunner system. Then, we detail how we format its output in such a that it fits the input of our truth discovering process.

Information extraction. We aim at analyzing conflicting claims about the same real-world facts from TextRunner. TextRunner [12, 2] is an OpenIE system relying on an unsupervised Web extraction process over a predefined set of Web corpus consisting of *Google*, *ClueWeb*, *News*, *Nell*, and *Wikipedia* corpus. Each particular corpus aggregates data from multiple other Web sources which are of various nature, for instance domain-specific Websites. Queries are sent to TextRunner by users through a form where they also can optionally specify their trusted corpus on the predefined set. When a user query arrives, TextRunner first finds the relevant sources in the set of corpus, and then it extracts each possible answer from them using natural language processing techniques and ontologies. To do so, it performs, for efficiency concerns, a single data-driven pass on the corpus to obtain the list of candidate relational tuples which might satisfy the arguments of the user input query. A typical user query in TextRunner is often about two real-world *entities* and a certain *relation* between them. As a consequence, such a user query q can be defined formally as a triplet (e_1, r, e_2) where e_1 and e_2 are real-world entities and r is a relation. The argument r specifies a possible relationship between the two given entities e_1 and e_2. In general, at least one among the three arguments is unkown, which captures a partial knowledge about the real world. In this context, TextRunner tries to find out the actual possible values of unkown arguments given known ones by querying the Web.

The outcome of TextRunner, given a user query, is indeed a set of candidate answers which is ranked according to the number of sources supporting each. TextRunner also enables to access, via Web hyperlinks, to the source and the document associated to each extracted answer. Our demonstration system (see Section 4) will follow these links and will extract the different sources of each potential answer

1. http://openie.allenai.org/

Figure 1: Data collection and formatting

for truth finding purposes.

Information processing. We consider a truth discovery process that takes, as input data, a set of claims in the form of quadruplets (claimID, source, claimQuery, claimValue) and infer, as outputs result, a *Boolean truth label* for each claim in which claimQuery is the query the claim is referred to and claimValue is the answer given by the source for the query. We detail below how these claims are inferred from the outcome of OpenIE TextRunner.

Assume a user query q about a real-word fact f_q and the set of n answers $v_1^q \ldots v_n^q$ returned by TextRunner for q. Let now denote by $\mathcal{S}_i^q$ the set of sources supporting each answer v_i^q, for $1 \leq i \leq n$. Recall that we extract this set of sources by following Web hyperlinks attached to answers by TextRunner and by using hand-written mapping rules. In addition note that for the same query, TextRunner returns only one answer per source, i.e. $\mathcal{S}_i^q \cap \mathcal{S}_j^q = \emptyset$ for all $i \neq j$ with $1 \leq i \leq j \leq n$. We have the set of potential answers along with their respective sets of sources for the user query q about the fact f_q. We can now proceed to the generation of the corresponding claims. To do so, we consider each answer v_i^q, for each $1 \leq i \leq n$, and loop on its set of sources $\mathcal{S}_i^q$. For each source $s \in \mathcal{S}_i^q$, we create a new claim (claimID, s, q, v_i^q) with an automatically generated unique claim identifier claimID, a claim source source corresponding to s, q as the claim query claimQuery, and a claim value claimValue equals to v_i^q. In other words, we create as many claims as the number of sources of a given answer from TextRunner. The total number of generated claims for the query q being equals to $\sum_{1 \leq i \leq n} |\mathcal{S}_i^q|$.

Let us assume the set of m successive user queries $q_1, \ldots, q_m$. We finally suppose that when the query is successively issued in TextRunner and answered by the system, the aforementioned formatting procedure transforms its set of possible answers to the corresponding set of claims. We refer to the overall set of resulting claims for queries $q_1, \ldots, q_m$ with $\mathcal{C}$.

[**Lamine: Types de requêtes á supporter?**]

[**Lamine: Traitement de requêtes en batch ou traitement une par une?**]

[**Lamine: Regarder la distribution du nombre de conflits par query sur TextRunner pour mieux motiver l'utilité du truth finding?**]

3. ENSEMBLING FOR TRUTH FINDING

We present in this section an adaptive truth finding approach which uses active ensembling in order to adaptively learn about an optimal set of truth finding algorithms that outperforms any individual technique on any given dataset. Our learning approach will actively involve users for the correct labels (or answers) of a sample of queries that cause maximal disagreement amongst our classifiers.

3.1 Ensemble-Based Active Learning

– donner idée générale pour introduire ce qu'est l'ensembling
– on a besoin de le faire dans le contexte de truth discov-

ery car aucune methode ne bat toutes les autres dans tous les cas de figure
– donc on combine les methodes : il y plusieurs façon de combiner par ex. consensus de méthodes, etc.
– expliquer quelles méthodes on combine avec leurs avantages et inconvénients

An ensemble-based active learning, or commonly ensembling, is a semi-supervised learning approach that tries to figure out an optimal ensemble of classifiers for a given classification problem by actively querying an oracle, e.g., a human being, about the labels of a sample of data items. Ensembling, thereby, enables to perform classification consistently well across datasets without having to determine a *priori* a suitable classifier type.

In our context, the truth finding algorithms correspond to our set of classifiers. The underlying *binary classification* problem consists of assigning the correct truth label to a set of claims about given user queries. Indeed, a truth finding algorithm is formally a mapping $\mathsf{TF} : \mathcal{C} \mapsto \{\mathsf{true}, \mathsf{false}\}$ which associates to each claim in $\mathcal{C}$ either true or false. A good truth finding algorithm provides predictions that are close to the actual world. Unfortunately, a well known property, e.g., as shown in [5, 9], of existing truth discovering algorithms remains their sentivity to certain application domains or datasets. As a consequence, there is no actual approach that outperforms the others on all types of datasets. On the other hand, truth finding is hard in practical scenarios because there is often no prior knowledge guiding to the selection, beforehand, of an optimal algorithm, in particular when the context is dynamic. More importantly, a large set of labeled examples (or ground truth) for evaluating the precisions of the algorithms is expensive to obtain in real applications.

In general, human being has a certain background knowledge about some real-world facts. Such a knowledge can serve as a valuable and inexpensive source of labels for a rather reasonable number of data items. However, having this partial ground truth from users is not sufficient in order to definitively decide about an optimal truth finding strategy because it can change over time as we obtain more information from sources, e.g. when claims are continuously extracted by TextRunner for answering new incoming queries. Therefore, there is a need for an adaptive approach able to dynamically figure out the optimal truth finding strategy when users' feedbacks and new knowledge about the world are available. We believe that active ensembling should be helpful to this end.

We put forward and demonstrate an approach which combines truth discovery and open information extraction with ensemble-based active learning for adaptively learning about the optimal ensemble of truth finding algorithms when the OpenIE system is gradually querying and labeled examples from users are available. As we shall show later, we will actively involve users to obtain the truth about a sample of particular facts during the learning process. The way this sampling is defined and chosen is crucial for the effectiveness

of the active learning. Several sample selection strategies, e.g., random sampling, query by committee, or support vector machine models, have been proposed for the definition of the type of selected data items along the size of the sample; we defer to [8] for more details about active machine learning. In this study, we use *query by committee* (QBC) for ensemble-based active learning. QBC states that the best data items to select for labels are those that cause the *maximal disagreement* among the predictions of an ensemble of diverse but partially accurate classifiers during active learning. Furthermore, we seek to provide an adaptive active learning by looking for an optimal ensemble given a larger set of input classifiers. [**Lamine:** Peut être qu'il y a mieux que QBC ?]

To learn about an optimal ensemble from a diverse set of classifiers, we have considered twelve well established truth finding algorithms in the literature, having three different types according to their specificities. Note that diversity offers better result in active learning than using homogenoeus classifiers (see [6]). We briefly present each considered class of truth discovering algorithms in the following.

1. **Iterative techniques:** TruthFinder [13], Cosine, 2-Estimates and 3-Estimates [3], AccuNoDep [1]
2. **EM based techniques:** MLE [11], LTM [14], SimpleLCA and GuessLCA [7]
3. **Dependency detection based techniques:** Depen, Accu, and AccuSim [1]

[**Lamine:** Peut être qu'il existe une meilleure classification ?]

3.2 Truth Finding with Active Ensembling

– notre approche que l'on défend ici dans la démo est semi supervisée en impliquant de l'utilisateur de façon active en lui demandant s'il peut confirmer des faits (facts)
– si on a une ground truth partielle on la "rejoue" cas par cas

We present an adaptive truth finding algorithm based on active ensembling in order to learn about an optimal ensemble over a set of existing truth finding algorithms, on which one can efficiently find the truth for the output of OpenIE systems. Our approach first obtains from the learning procedure intuitions about the best algorithm to use for each incoming fact (or a query about it) and then it performs the union of the result of the ensemble of best truth finding algorithms returned for a collection of facts. The best truth finding algorithm for a fact, i.e., the algorithm that has the highest chance to reliability discover the truth among a set of candidates claims about this fact, will be found by the learner by evaluating the accuracy of each competing technique on labeled claims from the user. We sketch in the following the procedure by assuming that all the input truth finding algorithms are used with their optimal initial parameters which are deemed known beforehand.

[**Lamine:** A sketch of the active learning process for truth finding] Our ensemble-based learning process relies on QBC for label querying and aims at finding an optimal combination of the results of the different truth finding algorithms involved in the process. Such a learning process is iterative and considers, as input, the set of unlabeled claims $\mathcal{C}$ and the set of labeled claims $\mathcal{C}_{GT}$. The claims in $\mathcal{C}$ are progressively obtained and processed, like in a streaming setting, from the information extraction system as it continuously receives queries. The set $\mathcal{C}_{GT}$, initially empty, contains claims whose labels are known for sure by asking the user. We also assume a stop criteria, e.g., a predefined number of iterations or accuracy changes between two iterations, for our iterative learning algorithm for truth finding and the set of optimal initial parameters for each truth finding algorithm. The algorithm starts by evaluating the truth finding algorithms on the available unlabeled claims in $\mathcal{C}$ for determining the prediction of each technique. It then compares those label predictions on each set of claims about the same query (thereby, an identical fact) and determines the queries causing the maximal disagreement among the members of the committee. Those queries are determined by computing the vote entropy of each query. The algorithm requests to the user labels for the set of candidates claims of those highly controversial queries. Once the labels are acquired from the user, the algorithm adds those labels along the associated claims into the labeled set $\mathcal{C}_{GT}$ and discards the claims from the unlabeled set $\mathcal{C}$. The learning procedure finally determines the accuracy of each input truth finding algorithm on the newly labeled claims in order to know the best technique for truth discovery.

1. The algorithm starts with an initialization phase in which values of the initial parameters of the learner and the different truth finding algorithm are set.
2. The algorithm pursues by giving the claims in $\mathcal{C}$ to the set of truth finding algorithms for label predictions and then it records the vote entropy of each query (thereby the underlying fact) according to the predictions of the committee.
3. It chooses the set of claims associated to the query having the maximum vote entropy for label querying
4. The set of acquired truth labels, together with the corresponding claims, are added into $\mathcal{C}_{GT}$ and then discarded from $\mathcal{C}$.
5. At this stage the algorithm estimates the accuracy of the different truth finding algorithm on $\mathcal{C}_{GT}$ in order to determine the best one for truth discover over those claims.

The steps 2–5 of the active learning algorithm are repeated until the stop criteria is satisfied. At the end of the active learning process, the truth discovery is finally realized by performing the union of the predictions of the ensemble of best truth finding algorithms for the collection of claims.

4. OUR DEMONSTRATION SYSTEM

We describe in this section our system for combining truth discovring and information extraction with active ensembling using the procedure described in the previous sections. We first present the architecture of our system by giving its different modules. Then, we provide a typical demonstration scenario of a user interacting with our system.

4.1 System Components

The architecture of our demonstration system, given in Figure 2, comprises the following five main components.

User Interface. The user interface is the main component that enables a given user to interact with our system. Via the user interface component, one has the ability to provide a query (or a key phrase), to recieve answers from the truth

finding module or label requests from the active ensembling module. When a label request is sent is sent to the user, she (or he) also provides his answers through this component.

OpenIE component. OpenIE is responsible to the extraction of information from the Web. It considers, as an input, a user query and queries several Web corpus in order to returns the relevant set of candidate answers. Our OpenIE component relies on TextRunner engine for the extraction of the set of candidate Web claims with respect to a user query about a given real-world fact.

Active Ensembling Module. This is the core component of our system which discovers an optimal ensemble for truth finding over candidate claims extracted with the OpenIE component. The active ensembling module, as described in Section 3, evaluates set of twelves truth finding algorithms on unlabeled claims, figures out the most controversal claims, and requests labels from the user. Once it obtained feebacks from the user, it estimates the accuracy of the competing algorithms on the labeled claims, and finally decides about the best one to use for each set of claims related to the same query or fact. The module returns an ensemble since one can have distinct good truth finding techniques for various sets of claims about different facts.

Truth Finding Module. The truth finding module uses the hypotheses found by the active ensembling module about the best algorithm to use for each set input of claims to determine the final result that will maximize the precision of the truth finding discovery. Typically, the module combines the best of each algorithm in the ensemble returned by the learning process. It outputs its result to the user through the user interface component.

Storage Module. The storage module consists of a local repository for the system that enables to have a dump of labeled claims. This set of labeled claims is a partial valuable ground truth. It can be used, if available, in order to boostrap our active learning process. For example as we already learnt about the best approach for these labeled claims, they can be compare to new unlabeled claims, in terms of similar involved facts or similar conflict distribution, for directly devising a candidate optimal truth discovery algorithm without having to run again the learning procedure.


Figure 2: Architecture of our system

4.2 Demonstration scenario

A given user that wants to interact with our system must do it through the search form. Through the search form, she (or he) provides her searched relation, e.g., "Where is

born Barack Obama?". The searched relation is then passed to the information extraction engine, TextRunner system in our case, which returns a set of answers considered to be relevant for the user's request. Each claim in the returned list is processed in order to extract the corresponding sources along a detailed description of the claim which we format in a certain manner. The set of sources and the formatted versions of all claims are then passed to the truth finding module which integrate all the claims and compute the most probable answer together with the reliability scores of participated sources for the searched relation. Finally, the output of the truth finding process is returned to the user. The user can also want to review the output of our system by definitively validiting it or not through its knwoledge of the modeled world. For example when the system has totally wrong, it may be interesting to get such a kind of feedbacks from the user in order to change the used method, as there are many available with our system, and to enhance the process for the further search about the same world. The user gives feedbacks using the option buttons on the left-hand side of the outputted claims or the text form. The feebacks given by the user is saved in knwoledge bases within our system for further processes.

5. CONCLUSION

[**Lamine:** L'utilisateur peut faire une erreur sur l'étiquette de certains claims. Comment capturer ce phénoméne?] []

6. REFERENCES

[1] Xin Luna Dong, Laure Berti-Equille, and Divesh Srivastava. Integrating conflicting data: The role of source dependence. *PVLDB*, 2(1):550–561, 2009.

[2] Oren Etzioni, Michele Banko, Stephen Soderland, and Daniel S. Weld. Open Information Extraction from the Web. *Commun. ACM*, 51(12):68–74, December 2008.

[3] Alban Galland, Serge Abiteboul, Amélie Marian, and Pierre Senellart. Corroborating Information from Disagreeing Views. In *WSDM*, pages 131–140, 2010.

[4] Naeemul Hassan, Chengkai Li, and Mark Tremayne. Detecting check-worthy factual claims in presidential debates. In *Proc. CIKM*, pages 1835–1838. ACM, 2015.

[5] Xian Li, Xin Luna Dong, Kenneth Lyons, Weiyi Meng, and Divesh Srivastava. Truth finding on the deep web: Is the problem solved? *Proc. VLDB Endow.*, 6:97–108, December 2012.

[6] Zhenyu Lu, Xindong Wu, and J.C. Bongard. Active learning through adaptive heterogeneous ensembling. *IEEE Transactions on Knowledge and Data Engineering*, 27:368–381, Feb 2015.

[7] Jeff Pasternack and Dan Roth. Latent Credibility Analysis. In *WWW*, pages 1009–1020, 2013.

[8] Burr Settles. *Active Learning*. Number 114. Morgan & Claypool Publishers,, June 2012.

[9] Dalia Attia Waguih and Laure Berti-Equille. Truth discovery algorithms: An experimental evaluation. *CoRR*, abs/1409.6428, 2014.

[10] Dalia Attia Waguih, Naman Goel, Hossam M. Hammady, and Laure Berti-Equille. AllegatorTrack: combining and reporting results of truth discovery

from multi-source data. In *Proc. ICDE*, pages
1440–1443, 2015.

[11] Dong Wang, Lance M. Kaplan, Hieu Khac Le, and
Tarek F. Abdelzaher. On Truth Discovery in Social
Sensing: a Maximum Likelihood Estimation Approach.
In *IPSN*, pages 233–244, 2012.

[12] Alexander Yates, Michael Cafarella, Michele Banko,
Oren Etzioni, Matthew Broadhead, and Stephen
Soderland. TextRunner: Open Information Extraction
on the Web. In *Proc. NAACL*, pages 25–26.
Association for Computational Linguistics, 2007.

[13] Xiaoxin Yin, Jiawei Han, and Philip S. Yu. Truth
Discovery with Multiple Conflicting Information
Providers on the Web. *TKDE*, 20(6):796–808, 2008.

[14] Bo Zhao, Benjamin I. P. Rubinstein, Jim Gemmell,
and Jiawei Han. A Bayesian Approach to Discovering
Truth from Conflicting Sources for Data Integration.
PVLDB, 5(6):550–561, 2012.